

## 1 Notations

We consider the prediction task of a set of observations  $(y^0, \dots, y^T)$  given a set of input  $(u^0, \dots, u^T)$ . The model can be built on top of a arbitrary encoder model, denoted  $h$ , that will be trained by gradient descent. Through this encoder, inputs are mapped to latent vectors:

$$h : (u^0, \dots, u^T) \mapsto (\tilde{u}^0, \dots, \tilde{u}^T)$$

## 2 Model

Our model can be viewed as a RNN layer followed by a fully connected layer. At time step  $t$ ,

$$\begin{cases} y_{t+1} = \tanh(W_y x_{t+1}^L + b_y) \\ x_{t+1} = \tanh(W_{xx} x_t + W_{xu} \tilde{u}_{t+1} + b_x) \end{cases}$$

with  $x_0 \equiv 0$ .

Let's consider the weights of the last RNN and fully connected layers as  $\theta \equiv (W_{xx}, \textcolor{red}{W}_{xu}, b_x, W_y, b_y)$ . We introduce two sequences of random noises as i.i.d real valued random variables  $\epsilon$  and  $\eta$ , with covariance matrices  $\Sigma_y$  and  $\Sigma_x$ . We can now write our model in terms of two functions  $f$  and  $g$  as:

$$\begin{cases} y_t = f_\theta(x_t) + \epsilon_t & \text{observation model} \\ x_{t+1} = g_\theta(x_t, u_{t+1}) + \eta_{t+1} & \text{state model} \\ \tilde{u}_t = h(u_t) & \text{command model} \end{cases} \quad (1)$$

In the following section, we will focus on maximizing the joint log likelihood:

$$\log p_\theta(X_{0:T}, y_{0:T}, u_{0:T}) \quad (2)$$

## 3 Minimization

We can start by developing the log likelihood:

$$\begin{aligned}
\log p_\theta(X_{0:T}, y_{0:T}, u_{0:T}) &= \frac{1}{T} \log \left( p_\theta(x_0) p_\theta(y_0|x_0) \prod_{k=1}^T p_\theta(x_k|x_{k-1}, u_k) p_\theta(y_k|x_k) \right) \\
&= \frac{1}{T} \log p_\theta(x_0) \\
&\quad + \frac{1}{T} \sum_{k=1}^T \log p_\theta(x_k|x_{k-1}, u_k) + \frac{1}{T} \sum_{k=0}^T \log p_\theta(y_k|x_k) \\
&= \frac{1}{T} \log p_\theta(x_0) - \frac{1}{2} \log |\Sigma_x| - \frac{1}{2} \log |\Sigma_y| + Cst \\
&\quad - \frac{1}{2T} \sum_{k=1}^T (x_k - g_\theta(x_{k-1}, u_k))' \Sigma_x^{-1} (x_k - g_\theta(x_{k-1}, u_k)) \\
&\quad - \frac{1}{2T} \sum_{k=0}^T (y_k - f_\theta(x_k))' \Sigma_y^{-1} (y_k - f_\theta(x_k))
\end{aligned}$$

We aim at maximizing 2 by gradient descent, by leveraging fisher's identity:

$$\nabla \log p_\theta(x_{0:T}, y_{0:T}, u_{0:T}) = \mathbb{E}_\theta [\nabla \log p_\theta(x_{0:T}, y_{0:T}, u_{0:T}) | Y_{0:T}]$$

In order to approximate this expectation, we need to sample from the posterior distribution  $p_\theta(x|y)$ . In Section 4, we detail a Sequential Monte Carlo approach to this end. In Section 5, we describe the algorithm to train our model through gradient descent.

## 4 Sequential Monte Carlo Approach

### 4.1 Filter

In order to compute the conditional expectations in the previous expressions, we will iteratively sample trajectories  $\xi_{1:T}^i$  associated with weights  $\omega^i$  with respect to the density  $p_\theta(x|y)$ , using a sequential Monte Carlo particle filter.

At time step  $k = 0$ ,  $(\xi_0^l)_{l=1}^N$  are sampled independently from the first hidden state, and associated with sampling weights proportional to the observation density  $q_\theta$ :

$$\begin{aligned}
\xi_0^i &\sim \mathcal{N}(x_0, \Sigma_x) \\
\omega_0^i &\sim q_\theta(\xi_0^i)
\end{aligned}$$

At time step  $k + 1$ , we sample indices  $I$  of the particles to propagate, based on their previous weights. After propagation, particles weights are computed following the observation density function:

$$\begin{aligned}
\mathbb{P}(I_{k+1}^i = j) &= \omega_k^j \quad \forall 1 \leq j \leq N \\
\omega_{k+1}^i &\sim q_\theta(\xi_{k+1}^i)
\end{aligned}$$

## 4.2 Smoother

Using the poor man filter, we get  $N$  trajectories:

$$\xi_{1:k+1}^i = (\xi_{1:k}^{I_{k+1}^i}, \xi_{k+1}^i)$$

## 4.3 Approximation

We can now approximate this conditional expectation for any measurable bounded function  $h$ :

$$\mathbb{E}_\theta [h(x)|y_{1:T}] = \sum_{i=1}^N \omega_T^i h(\xi_{0:T}^i)$$

# 5 Gradient descent

## 5.1 Forward pass

During the forward pass, we generate a set of  $N$  particles under the law  $p(x|y)$  for fixed values of  $\theta$ ,  $\Sigma_x$  and  $\Sigma_y$ . We initialize the sequence with a initial hidden state sampled from the noise's distribution.

$$x_0^i \sim \mathcal{N}(0, \Sigma_x)$$

In order to predict each new time step  $k + 1$ , particles from the previous step are attributed weights  $\omega_k^i$  proportionally to the density probability around the targeted value  $y_k$ .

$$\omega_k^i \sim \exp\left(-\frac{1}{2}(y_k - f_{\theta_p}(x_k^i))' \Sigma_{y,p}^{-1}(y_k - f_{\theta_p}(x_k^i))\right)$$

We then select a new population from these particles indexed by  $I_{k+1}^i$ , based on their weights.

$$\mathbb{P}(I_{k+1}^i = j) = \omega_k^j \quad \forall 1 \leq j \leq N$$

The current hidden state is computed for the selected particles.

$$x_{k+1}^i = g_\theta(x_k^{I_{k+1}^i}, u_{k+1}) + \eta_{k+1}^i$$

## 5.2 Loss function

Considering that we have computed a set of  $N$  trajectories  $(\xi_{0:T}^i)$ ,  $1 \leq i \leq N$ , associated with weights  $(\omega^i)$ , we can approximate the gradient of the log likelihood by computing the gradient of:

$$\begin{aligned} \mathbb{J}(\theta) &= \log |\Sigma_x| + \log |\Sigma_y| \\ &+ \frac{1}{T} \sum_{k=1}^T \sum_{i=1}^N \omega^i (y_k - f_\theta(\xi_k^i))' \Sigma_y^{-1} (y_k - f_\theta(\xi_k^i)) \\ &+ \frac{1}{T} \sum_{k=0}^T \sum_{i=1}^N \omega^i (\xi_k^i - g_\theta(\xi_{k-1}^i, u_k))' \Sigma_x^{-1} (\xi_k^i - g_\theta(\xi_{k-1}^i, u_k)) \end{aligned}$$

### 5.3 Backward pass

During this step, each parameter of the model is updated by gradient descent.

## 6 Two-step training

In this section, we aim at improving the weights of a model already trained in a traditional fashion.

### 6.1 Initial training

Given a dataset of samples  $(u_{0:T}^{(i)}, y_{0:T}^{(i)})_{i=1}^m$ , we consider a training of our model as defined in 1 without the added noise. For each sample, we simply compute a prediction  $\hat{y}_{0:T}$  by running the input through each layer, and minimize the discrepancy to the target values by gradient descent.

$$\mathbb{J} = \|u_{0:T} - y_{0:T}\|^2$$

This training results in an initial set of weights  $\theta_0$ .

### 6.2 Finetuning

Because we trust the initial training has successfully extracted relevant information from the command  $u$ , and reached satisfactory parameters for the hidden state model, we will only adapt the parameters of the observation model:

$$\theta_{finetune} \equiv (\cancel{W_{xx}}, \cancel{b_x}, W_y, b_y)$$

$\Sigma_x$  is set to a small value and will not be updated during the finetuning.